

并行与分布式计算基础：第二十一讲

杨超

chao_yang@pku.edu.cn

2019 秋



内容提纲

1 CUDA 编程-5

- 补遗

补遗


1 CUDA 编程-5

- 补遗

线程簇分歧

- 回顾：线程块将按照线程簇 (一般为 32 个线程) 为单元在 SM 上调度，同一线程簇中所有线程采用 SIMD (或称 SIMT) 方式执行。
- 线程簇分歧 (warp divergence)：当同一线程簇中线程执行不同程序路径时，会触发串行执行，导致程序性能下降。

```
if (...) {  
    // Path A  
} else {  
    // Path B  
}
```



例如，在向量加法中的边界检查：

```
if (i < n)d_C[i] = d_A[i] + d_B[i];
```

- 向量长度为 100 时，总共 4 个 warp，其中有 1 个 warp 产生分歧，占 25%.
- 向量长度为 1000 时，总共 32 个 warp，其中有 1 个 warp 产生分歧，占 3%.


可以看出，对于边界检查，一般来说随着规模增加，分歧的影响会降低.

常见分歧的处理策略

- 线程分支：比如，在代码 `if (threadIdx.x > 2)...` 中，线程 0, 1, 2 和线程 3-31 执行路径不同，程序分支尽量以线程簇大小作为粒度，设法改为 `if (threadIdx.x / WARP_SIZE < 2)...`；
- 边界检查：例如向量加法中 `if (i < n)d_C[i] = d_A[i] + d_B[i];`，如果开销太大，可以考虑使用两个 kernel，一个处理边界内的计算，一个处理边界情况。
- 线程改变：一些并行算法如 reduction 等，随着时间推移，参与的线程数目发生改变，可以通过设计新的并行算法来减少线程簇分歧。

线程簇混洗

- 线程簇混洗 (warp shuffle): 允许某线程直接读取同一个 warp 中其它线程寄存器中的值, 延迟低且不占用额外的存储资源.



原子操作


- 原子操作 (atomic operations): 当所有线程同时修改某个全局变量时, 需要加锁后访问, 保证结果的正确性, 常见的原子操作有:

```
atomicAdd, atomicSub, atomicMin, atomicMax,  
atomicInc, atomicDec, atomicExch, atomicCAS,  
atomicAnd, atomicOr, atomicXor
```

- 线程簇聚合 (warp aggregation): 多个线程原子累加到单个计数器以提高性能, 线程簇中的线程首先计算它们之间的总增量, 然后选择单个线程将增量原子地添加到全局计数器中;
- CUDA9.0 以上的 NVCC 编译器已官方支持自动的为原子操作执行线程簇聚合.


协同分组

- 协同分组 (cooperative groups): CUDA9.0 以上, 支持更为灵活的线程组合方式, 从而可以在不同粒度上进行线程间协作.



向量化访存

- 向量化访存 (vectorized memory access): 对于访存受限的问题, 在保证数据对其的前提下, 通过使用例如 `float2` 等向量化的数据类型, 并结合 `reinterpret_cast` 对指针进行强制转换, 可以帮助编译器实现访存的向量化, 从而提升性能.




NVIDIA GPU 的计算能力 (compute capability)

- 计算能力用于反映 CUDA 设备所支持的不迭更迭的功能和特性；
- 计算能力以 x.y 表示，两个数字分别为主版本和从版本号；
- 计算能力的版本之间向后兼容，越新表示设备的功能越强大；




Kepler
CC 3.5
192 cores / SMX



Maxwell
CC 5.0
128 cores / SMM



Pascal
CC 6.0
64 cores / SMM



Volta
CC 7.0
64 cores / SMM

<https://developer.nvidia.com/cuda-gpus>

12 NVIDIA

- 不同计算能力、不同产品线的设备构成了庞大的生态体系。

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT, cuBLAS, cuRAND, cuSPARSE	CULA MAGMA	Thrust NPP	VSIP, SVM, OpenCurrent	PhysX, OptiX, iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java, Python, Wrappers	DirectCompute	Directives (e.g., OpenACC)	
CUDA-enabled NVIDIA GPUs						
Turing Architecture (Compute capabilities 7.x)		DRIVE/JETSON AGX Xavier	GeForce 2000 Series	Quadro RTX Series	Tesla T Series	
Volta Architecture (Compute capabilities 7.x)		DRIVE/JETSON AGX Xavier			Tesla V Series	
Pascal Architecture (Compute capabilities 6.x)		Tegra X2	GeForce 1000 Series	Quadro P Series	Tesla P Series	
Maxwell Architecture (Compute capabilities 5.x)		Tegra X1	GeForce 900 Series	Quadro M Series	Tesla M Series	
Kepler Architecture (Compute capabilities 3.x)		Tegra K1	GeForce 700 Series GeForce 600 Series	Quadro K Series	Tesla K Series	
		EMBEDDED	CONSUMER DESKTOP, LAPTOP	PROFESSIONAL WORKSTATION	DATA CENTER	

- CUDA 提供了 deviceQuery 样例程序用于检查设备的计算能力。

```

$ ./deviceQuery
./deviceQuery Starting...

CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "Quadro GV100"
  CUDA Driver Version / Runtime Version      10.1 / 10.1
  CUDA Capability Major/Minor version number: 7.0
  Total amount of global memory:              32508 MBytes (34087305216 bytes)
  (80) Multiprocessors, ( 64) CUDA Cores/MP: 5120 CUDA Cores
  GPU Max Clock rate:                        1627 MHz (1.63 GHz)
  Memory Clock rate:                         850 Mhz
  Memory Bus Width:                          4096-bit
  L2 Cache Size:                             6291456 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
  Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total number of registers available per block: 65536
  Warp size:                                 32
  Maximum number of threads per multiprocessor: 2048
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
  Maximum memory pitch:                      2147483647 bytes
  Texture alignment:                          512 bytes
  Concurrent copy and kernel execution:       Yes with 4 copy engine(s)

```

Feature Support	Compute Capability					
(Unlisted features are supported for all compute capabilities)	3.0	3.2	3.5, 3.7, 5.0, 5.2	5.3	6.x	7.x
Atomic functions operating on 32-bit integer values in global memory (Atomic Functions)	Yes					
atomicExch() operating on 32-bit floating point values in global memory (atomicExch())	Yes					
Atomic functions operating on 32-bit integer values in shared memory (Atomic Functions)	Yes					
atomicExch() operating on 32-bit floating point values in shared memory (atomicExch())	Yes					
Atomic functions operating on 64-bit integer values in global memory (Atomic Functions)	Yes					
Atomic functions operating on 64-bit integer values in shared memory (Atomic Functions)	Yes					
Atomic addition operating on 32-bit floating point values in global and shared memory (atomicAdd())	Yes					
Atomic addition operating on 64-bit floating point values in global memory and shared memory (atomicAdd())	No				Yes	
Warp vote and ballot functions (Warp Vote Functions)	Yes					
__threadfence_system() (Memory Fence Functions)						
__syncthreads_count(),						
__syncthreads_and(),						
__syncthreads_or() (Synchronization Functions)						
Surface functions (Surface Functions)						
3D grid of thread blocks						
Unified Memory Programming						
Funnel shift (see reference manual)	No	Yes				
Dynamic Parallelism	No		Yes			
Half-precision floating-point operations: addition, subtraction, multiplication, comparison, warp shuffle functions, conversion	No			Yes		
Tensor Core	No					Yes

	Compute Capability											
Technical Specifications	3.0	3.2	3.5	3.7	5.0	5.2	5.3	6.0	6.1	6.2	7.0	7.5
Maximum number of resident grids per device (Concurrent Kernel Execution)	16	4	32				16	128	32	16	128	
Maximum dimensionality of grid of thread blocks	3											
Maximum x-dimension of a grid of thread blocks	2 ³¹ -1											
Maximum y- or z-dimension of a grid of thread blocks	65535											
Maximum dimensionality of thread block	3											
Maximum x- or y-dimension of a block	1024											
Maximum z-dimension of a block	64											
Maximum number of threads per block	1024											
Warp size	32											
Maximum number of resident blocks per multiprocessor	16				32						16	
Maximum number of resident warps per multiprocessor	64											32
Maximum number of resident threads per multiprocessor	2048											1024
Number of 32-bit registers per multiprocessor	64 K			128 K		64 K						
Maximum number of 32-bit registers per thread block	64 K	32 K	64 K				32 K	64 K		32 K	64 K	
Maximum number of 32-bit registers per thread	63	255										
Maximum amount of shared memory per multiprocessor	48 KB			112 KB	64 KB	96 KB	64 KB		96 KB	64 KB	96 KB	64 KB
Maximum amount of shared memory per thread block ²⁷	48 KB										96 KB	64 KB
Number of shared memory banks	32											
Amount of local memory per thread	512 KB											
Constant memory size	64 KB											
Cache working set per multiprocessor for constant memory	8 KB							4 KB	8 KB			

不同计算能力的 GPU 所支持的最大资源数

GPU	Kepler GK110	Maxwell GM200	Pascal GP100
Compute Capability	3.5	5.2	6.0
Threads / Warp	32	32	32
Max Warps / Multiprocessor	64	64	64
Max Threads / Multiprocessor	2048	2048	2048
Max Thread Blocks / Multiprocessor	16	32	32
Max 32-bit Registers / SM	65536	65536	65536
Max Registers / Block	65536	32768	65536
Max Registers / Thread	255	255	255
Max Thread Block Size	1024	1024	1024
Shared Memory Size / SM	16 KB/32 KB/48 KB	96 KB	64 KB

占有率

- 占有率 (occupancy): 每个 SM 中活跃线程簇数与最大线程簇数的比值, 越接近 100% 一般越好.
- 实际占有率往往受限于 kernel 的硬件资源消耗, 主要硬件资源有:
 - ▶ 每个 SM 寄存器的容量;
 - ▶ 每个 SM 共享内存的容量;
 - ▶ 每个 SM 允许的最大线程块数;
 - ▶ 每个 SM 允许的最大线程数;
 - ▶ 每个线程块允许的最大线程数.
- 上述因素与具体硬件的计算能力密切相关.

几个例子

P100 每个 SM 寄存器容量为 256KB、共享内存容量为 64KB、最大活跃线程簇数 64 (即 2048 线程)，最大活跃线程块数 32。

例子 1：寄存器限制

若每个线程使用 32 个单精度寄存器，则活跃线程为 $256\text{KB} \div (4\text{B} \times 32) = 2048$ ，占有率为 1；若每个线程使用 64 个单精度寄存器，占有率则降低为 0.5。

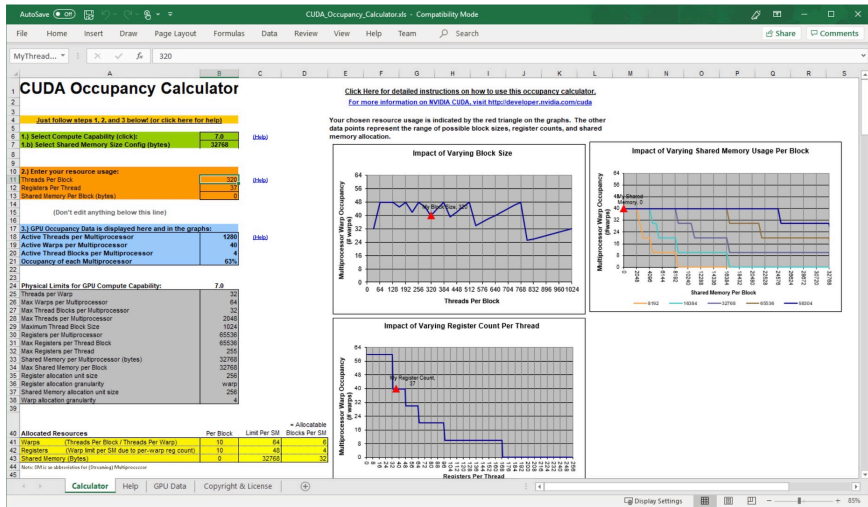
例子 2：共享内存限制

若每个线程使用 32Byte 共享内存，则活跃线程为 $64\text{KB} / 32\text{B} = 2048$ ，占有率为 1；若每个线程使用 64Byte 共享内存，占有率则降低为 0.5。

例子 3：线程块大小限制

若每个线程块有 32 个线程，占有率则不会高于 $32 \times 32 \div 2048 = 0.5$ 。

CUDA 占有率计算工具



CUDA 占有率计算函数

- 根据 kernel 预估块大小:

```
1  CUresult cuOccupancyMaxPotentialBlockSize(  
2      int* minGridSize, int* blockSize, CUfunction func,  
3      CUoccupancyB2DSize blockSizeToDynamicSMemSize,  
4      size_t dynamicSMemSize, int  blockSizeLimit )
```

- 返回 kernel 实际运行块数:

```
1  CUresult cuOccupancyMaxActiveBlocksPerMultiprocessor(  
2      int* numBlocks, CUfunction func,  
3      int blockSize, size_t dynamicSMemSize )
```